

Name : Harshal Kodgire
PRN : 2019BTECS00029
Batch : B1
Sub : HPC Lab

Assignment - 9

Q1) Implement Vector-Vector addition using CUDA C. State and justify the speedup using different sizes of threads and blocks.

Code :

```
#include <stdio.h>

void initWith(float num, float *a, int N)
{
    for(int i = 0; i < N; ++i)
    {
        a[i] = num;
    }
}

/*
 * Device kernel stores into `result` the sum of each
 * same-indexed value of `a` and `b`.
 */

__global__ void addVectorsInto(float *result, float *a, float *b, int
N)
{
    int index = threadIdx.x + blockIdx.x * blockDim.x;
    int stride = blockDim.x * gridDim.x;

    for(int i = index; i < N; i += stride)
    {
        result[i] = a[i] + b[i];
    }
}

void checkElementsAre(float target, float *vector, int N)
```

```

{
    for(int i = 0; i < N; i++)
    {
        if(vector[i] != target)
        {
            printf("FAIL: vector[%d] - %0.0f does not equal %0.0f\n", i,
vector[i], target);
            exit(1);
        }
    }
    printf("Success! All values calculated correctly.\n");
}

int main()
{
    const int N = 2<<24;
    size_t size = N * sizeof(float);

    float *a;
    float *b;
    float *c;

    cudaMallocManaged(&a, size);
    cudaMallocManaged(&b, size);
    cudaMallocManaged(&c, size);

    addVectorsInto<<<8,32>>>(c,a,b,N);
    cudaDeviceSynchronize();

    cudaFree(a);
    cudaFree(b);
    cudaFree(c);
}

```

For <<<1,1>>>

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
89.0	2197177162	1	2197177162.0	2197177162	2197177162	cudaDeviceSynchronize
10.6	261585252	3	87195084.0	23727	261490522	cudaMallocManaged
0.5	11147162	3	3715720.7	3650969	3813264	cudaFree
0.0	39971	1	39971.0	39971	39971	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	2197163231	1	2197163231.0	2197163231	2197163231	addVectorsInto(float*, float*, float*, int)

For <<<2,32>>>

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
55.2	305749237	1	305749237.0	305749237	305749237	cudaDeviceSynchronize
42.6	235994217	3	78664739.0	16791	235930900	cudaMallocManaged
2.3	12597439	3	4199146.3	4141923	4274340	cudaFree
0.0	28649	1	28649.0	28649	28649	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	305733663	1	305733663.0	305733663	305733663	addVectorsInto(float*, float*, float*, int)

For <<<4,64>>>

CUDA API Statistics:						
Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
64.7	242791066	3	80930355.3	17296	242726305	cudaMallocManaged
30.5	114271587	1	114271587.0	114271587	114271587	cudaDeviceSynchronize
4.8	17923683	3	5974561.0	5936920	6045652	cudaFree
0.0	33458	1	33458.0	33458	33458	cudaLaunchKernel
CUDA Kernel Statistics:						
Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	114256989	1	114256989.0	114256989	114256989	addVectorsInto(float*, float*, float*, int)

Q2) Implement N-Body Simulator using CUDA C. State and justify the speedup using different sizes of threads and blocks.

Code :

```
#include <math.h>
#include <stdio.h>
#include <stdlib.h>
#include "timer.h"
#include "files.h"

#define SOFTENING 1e-9f

/*
 * Each body contains x, y, and z coordinate positions,
 * as well as velocities in the x, y, and z directions.
 */

typedef struct { float x, y, z, vx, vy, vz; } Body;

/*
 * Calculate the gravitational impact of all bodies in the system
 * on all others.
 */

__global__ void bodyForce(Body *p, float dt, int n) {

    int i = blockIdx.x * blockDim.x + threadIdx.x;
    //int stride = gridDim.x * blockDim.x;
    if( i < n){

        float Fx = 0.0f; float Fy = 0.0f; float Fz = 0.0f;

        for (int j = 0; j < n; j++) {
            float dx = p[j].x - p[i].x;
            float dy = p[j].y - p[i].y;
            float dz = p[j].z - p[i].z;
            float distSqr = dx*dx + dy*dy + dz*dz + SOFTENING;
            float invDist = rsqrtf(distSqr);
            float invDist3 = invDist * invDist * invDist;
```

```

        Fx += dx * invDist3; Fy += dy * invDist3; Fz += dz *
invDist3;
    }

    p[i].vx += dt*Fx; p[i].vy += dt*Fy; p[i].vz += dt*Fz;

}
}

int main(const int argc, const char** argv) {

    // The assessment will test against both 2<<11 and 2<<15.
    // Feel free to pass the command line argument 15 when you generate
./nbody report files
    int nBodies = 2<<11;
    if (argc > 1) nBodies = 2<<atoi(argv[1]);

    // The assessment will pass hidden initialized values to check for
correctness.
    // You should not make changes to these files, or else the assessment
will not work.
    const char * initialized_values;
    const char * solution_values;

    if (nBodies == 2<<11) {
        initialized_values = "09-nbody/files/initialized_4096";
        solution_values = "09-nbody/files/solution_4096";
    } else { // nBodies == 2<<15
        initialized_values = "09-nbody/files/initialized_65536";
        solution_values = "09-nbody/files/solution_65536";
    }

    if (argc > 2) initialized_values = argv[2];
    if (argc > 3) solution_values = argv[3];

    const float dt = 0.01f; // Time step
    const int nIters = 10; // Simulation iterations

    int bytes = nBodies * sizeof(Body);
    float *buf;

```

```

cudaMallocManaged(&buf, bytes);

Body *p = (Body*)buf;

read_values_from_file(initialized_values, buf, bytes);

double totalTime = 0.0;

/*
 * This simulation will run for 10 cycles of time, calculating
gravitational
 * interaction amongst bodies, and adjusting their positions to
reflect.
 */

for (int iter = 0; iter < nIters; iter++) {
    StartTimer();

    /*
     * You will likely wish to refactor the work being done in
`bodyForce`,
     * and potentially the work to integrate the positions.
     */
    size_t threads_per_block = 256;
    size_t number_of_blocks = (nBodies + threads_per_block - 1) /
threads_per_block;

    bodyForce<<<number_of_blocks, threads_per_block>>>(p, dt, nBodies);
// compute interbody forces
    cudaDeviceSynchronize();

    /*
     * This position integration cannot occur until this round of
`bodyForce` has completed.
     * Also, the next round of `bodyForce` cannot begin until the
integration is complete.
     */

    for (int i = 0 ; i < nBodies; i++) { // integrate position
        p[i].x += p[i].vx*dt;
        p[i].y += p[i].vy*dt;
        p[i].z += p[i].vz*dt;
    }
}

```

```

    const double tElapsed = GetTimer() / 1000.0;
    totalTime += tElapsed;
}

double avgTime = totalTime / (double)(nIters);
float billionsOfOpsPerSecond = 1e-9 * nBodies * nBodies / avgTime;
write_values_to_file(solution_values, buf, bytes);

// You will likely enjoy watching this value grow as you accelerate
the application,
// but beware that a failure to correctly synchronize the device
might result in
// unrealistically high values.
printf("%0.3f Billion Interactions / second\n",
billionsOfOpsPerSecond);

cudaFree(buf);
}

```

In [16]: `run_assessment()`

Running nbody simulator with 4096 bodies

Application should run faster than 0.9s

Your application ran in: 0.2437s

Your application reports 15.895 Billion Interactions / second

Your results are correct

Running nbody simulator with 65536 bodies

Application should run faster than 1.3s

Your application ran in: 0.5703s

Your application reports 110.704 Billion Interactions / second

Your results are correct

Congratulations! You passed the assessment!

See instructions below to generate a certificate, and see if you can accelerate the simulator even more!